

# Adaptive Query Algorithms for Efficiently Discovering Overconfident Unknowns

Karsten Maurer

November 12th, 2019

# Who am I?



## **Karsten Maurer**

- Minnesota, Morris for BA in History and Statistics
- Iowa State for Masters and PhD in Statistics
- Currently an Assistant Professor of Statistics at Miami University
- Teaching: Data Visualization, Predictive Modeling, Statistical Programming
- Research: Statistics Education, Machine Learning, Visualization, Collaborative Consulting

# Overview of Today's Talk

## Unknown Unknowns

- Existing Literature on Utility-Based Query Algorithms
- Deficiencies in Existing Methods

## Facility Locations Search Algorithm

- Adapting Utility Objectives to Overcome Deficiencies
- Facility Locations Utility Function
- Greedy Adaptive Query Algorithm
- Empirical Experiment Results

## Adversarial Distances for Unsupervised Query Set Selection

- Objective and Intuition
- Adversarial Distance-Based Query Algorithm
- Empirical Experiment Results

## Discussion

- Conclusions, Current Work and Future Work

# Unknown Unknowns

## Unknown Unknowns

**Context:** You have a classifier that you wish to apply to a domain in which you only have access to a large set of unlabeled test instances.

**Problem:** You know that the classifier was trained in a different domain, or you suspect a bias in the training set relative to the testing

**Goal:** Select a set of instances from the unlabeled set that we will look up true labels which we use to evaluate the classifier performance. This is called a *query* of the unlabeled set.

An *oracle* is the person who will find the true labels. We don't want to waste the oracle's time, so we need a smart way to query the unlabeled set.

# Unknown Unknowns

## What are we looking for in the query set?

Attenberg et al. (2015) define *Unknown Unknowns* (UUs) as instances,  $x$ , from unlabeled test set where the classifier,  $f(\cdot)$ , is both:

- Highly confident. Above some threshold  $\tau \in (0, 1)$ . Thus,  $C_x \geq \tau$
- Wrong. Misclassified with  $y \neq f(x)$

## Why would we care about finding UUs in query set?

- High confidence mistakes lead to unmitigated risks in application of classifier
- Characteristics of the UUs may help analyst to understand classifier deficiencies

# Unknown Unknowns

## **Example: Image Classifier for Cats and Dogs**

- Bias in training set: Cats with light fur and Dogs with dark fur
- Test set has all color fur for both cats and dogs
- Classifier likely to be highly confident but wrong about dark cats and light dogs
- A query set that finds these UUs could help us to identify the model deficiency from animal fur color

# Existing Literature on Utility-Based Query Algorithms

## **Attenberg et al. (2015) - Crowdsourcing**

- "Beat the Machine" game with monetary rewards for finding UUs
- Crowdsourcing as mechanism for learning classifier deficiencies
- Basically like using many oracles with large budget for labelling

## **Lakkaraju et al. (2017) - Multi-Armed Bandit**

- Adaptive query algorithm that updates the optimal recommendation for next oracle query after evaluating the newly labeled instance
- Utility function provides a unit value for each discovered UU and penalizes by the cost of labeling

## **Bansal and Weld (2018) - Coverage-Based**

- Another greedy query algorithm, but uses coverage-based utility
- Goal is to encourage both discovery of UUs and exploration of feature space



## Coverage-Based Utility Details

$$U(Q) = \sum_{x \in \mathbb{X}} c_x \cdot \max_{q \in S} \{sim(x, q)\}$$

- $\mathbb{X} \subset \mathbb{R}^p$  is  $p$ -dimensional unlabeled test set
- $Q \subset \mathbb{X}$  is the set of points labeled by an oracle,
- $S = \{x | x \in Q, y_x \neq M(x)\}$  is the set of discovered UUs
- Classifier  $M(x) : \mathbb{X} \rightarrow class$ ,
- $c_x$  is the classifier's confidence in its prediction of  $x$ ,
- $sim(x, q)$  is a distance-based similarity metric.

## Deficiencies in existing methods

### Attenberg et al. (2015)

- If the cost of oracle queries is the reason we need an algorithm, then crowd-sourcing the solution is not viable

### Lakkaraju et al. (2017)

- Fundamentally place value on finding *any* UU, regardless of confidence
- This ignores the number of misclassification we **expect** based on the confidence scores

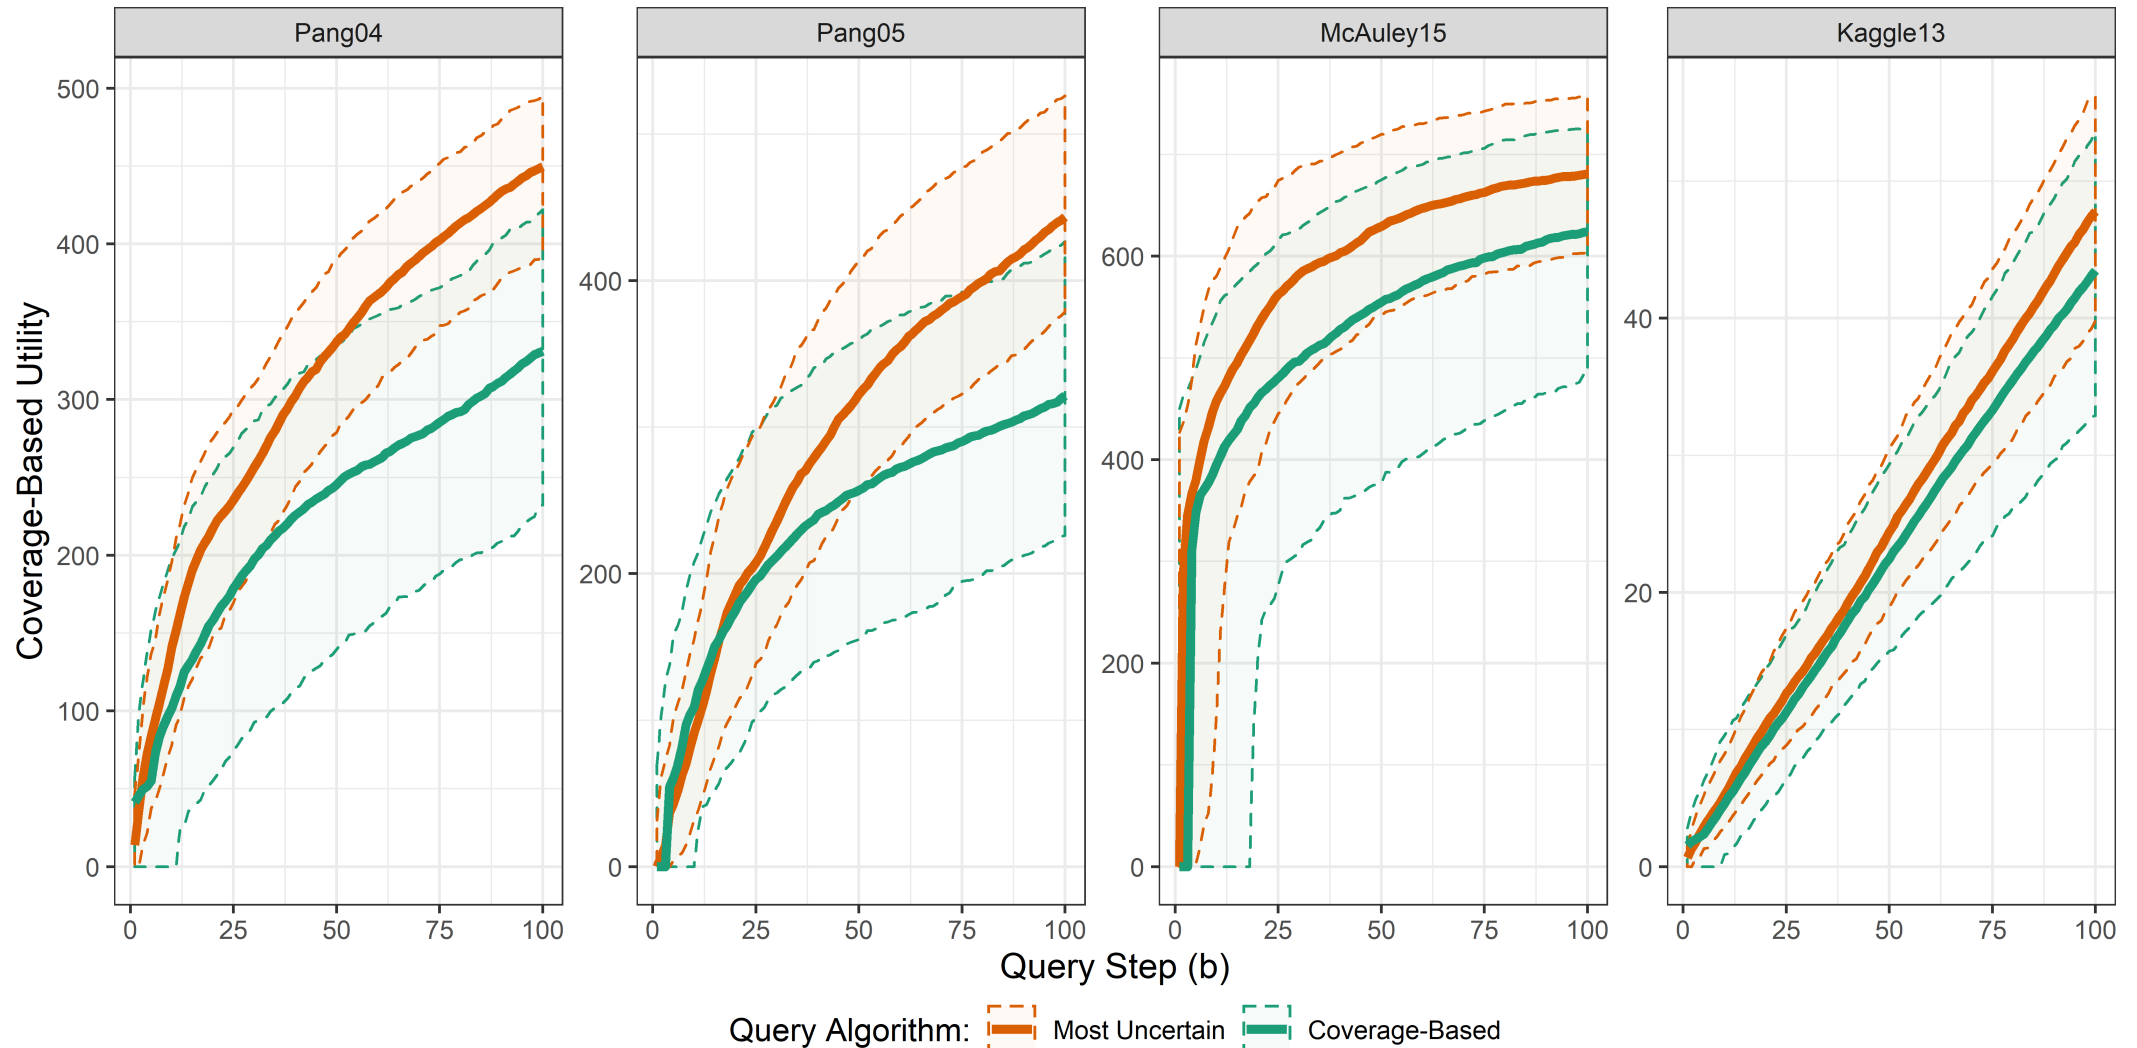
### Bansal and Weld (2018)

- Places value on high confidence instances being *close to* UUs

## Why is this a problem?

- The utility-based algorithms err toward instances that are "safe bets"; instances with confidence just above the  $\tau$  threshold.
- The utility functions reward selection of the most-uncertain (least confident) UUs

# Coverage-Based Query vs. Most-Uncertain Query



# Facility Locations Search Algorithm

# Adapting utility objectives to overcome deficiencies

## Claims:

The focus on discovering Unknown Unknowns, as defined in current literature, is flawed.

High confidence should not be interpreted as an absolute. Misclassifications *should* be occurring.

The problem is when the rate of misclassification exceeds the rate expected based on classifier confidence.

## Adaptated Objectives:

1. Our goal should be to query instances that demonstrate *classifier overconfidence*; containing more misclassification than should be expected based on classifier confidence.
2. Our query should also seek to thoroughly explore the feature space.

Thus, we propose a new utility function and associated greedy query algorithm that rewards both of these new objectives

## Facility Locations Utility Function

$$W(Q) = \sum_{q \in S} r(c_q) - \frac{1}{n} \sum_{x \in \mathbb{X}} \min_{q \in S} (d(x, q))$$

$r(c_q) = \log(1/(1 - c_q))$  is the *reward* function for finding a misclassification with confidence  $c_q$

$d(x, q)$  is the Euclidean distance between points  $x$  and  $q$ .

# Algorithm for Greedy Facility Location Search

## Input:

Test set  $\mathbb{X}$ , prior  $\hat{\phi}(x|Q = \emptyset)$ , budget  $B$

$Q = \{\}$  inputs that have been queried

$y_Q = \{\}$  oracle defined labels

**For:**  $b = 1, 2, \dots, B$  **do:**

$$q' = \operatorname{argmax}_{q' \notin Q} E_{\hat{\phi}}[W(Q \cup q')]$$

$$Q \leftarrow Q \cup q'$$

$$y_Q \leftarrow y_Q \cup y_{q'}$$

$$S \leftarrow \{x | x \in Q \text{ and } y_x \neq M(x)\}$$


$$\hat{\phi} \leftarrow \hat{\phi}(x|Q)$$

$$b \leftarrow b + 1$$

**Return:**  $Q, S$  and  $y_Q$

# Empirical Experiment Results

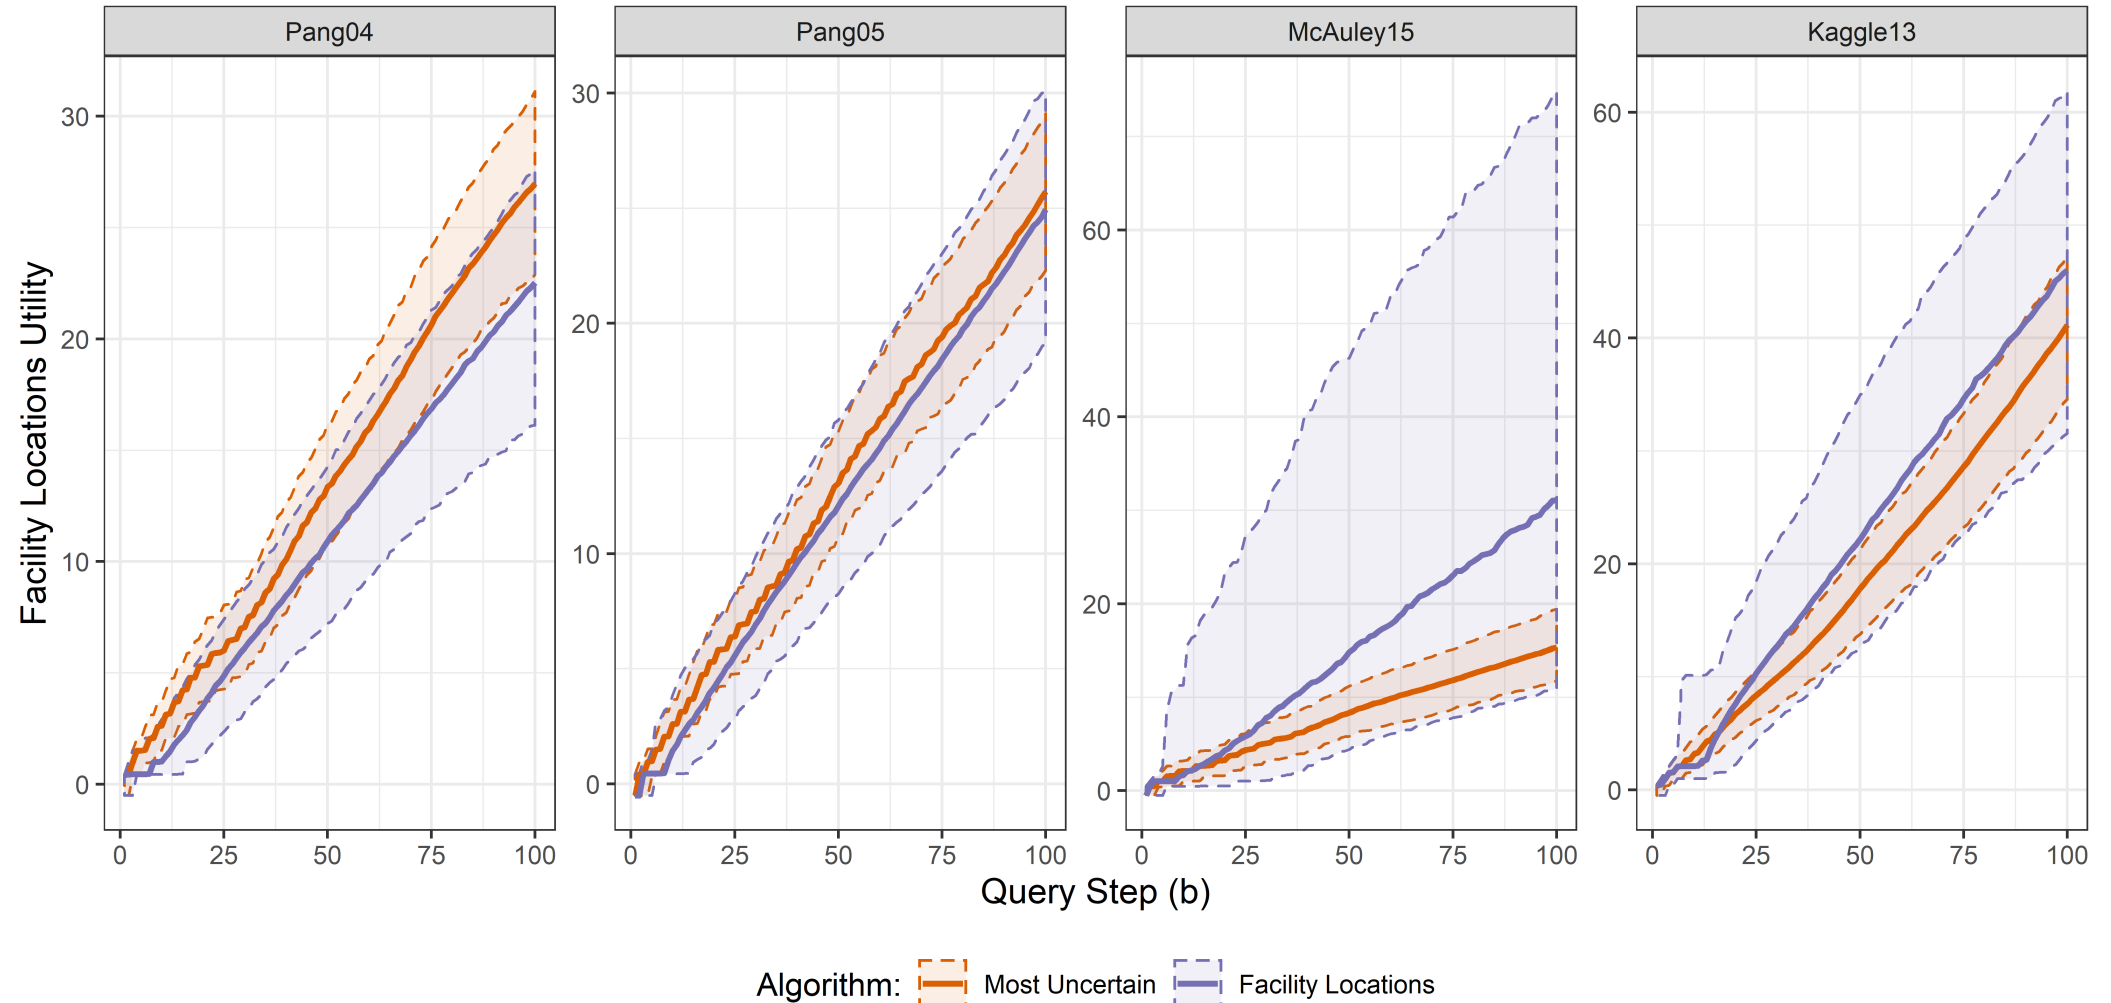
## Datasets



1. Compare Facility Location Query vs. Most-Uncertain Query baseline
2. Evaluate the efficiency of query algorithms in finding misclassification relative to the confidence of the selected instances



# Facility Location Query vs. Most-Uncertain Query



## Comparing Apples and Oranges

Making comparison to a baseline on different utility outcomes is too indirect

Can't use the utilities on which the query algorithms optimize without obvious bias

**Need a metric for comparison of query algorithms.**

$$|S| / \sum_{x=1}^B (1 - c_x)$$

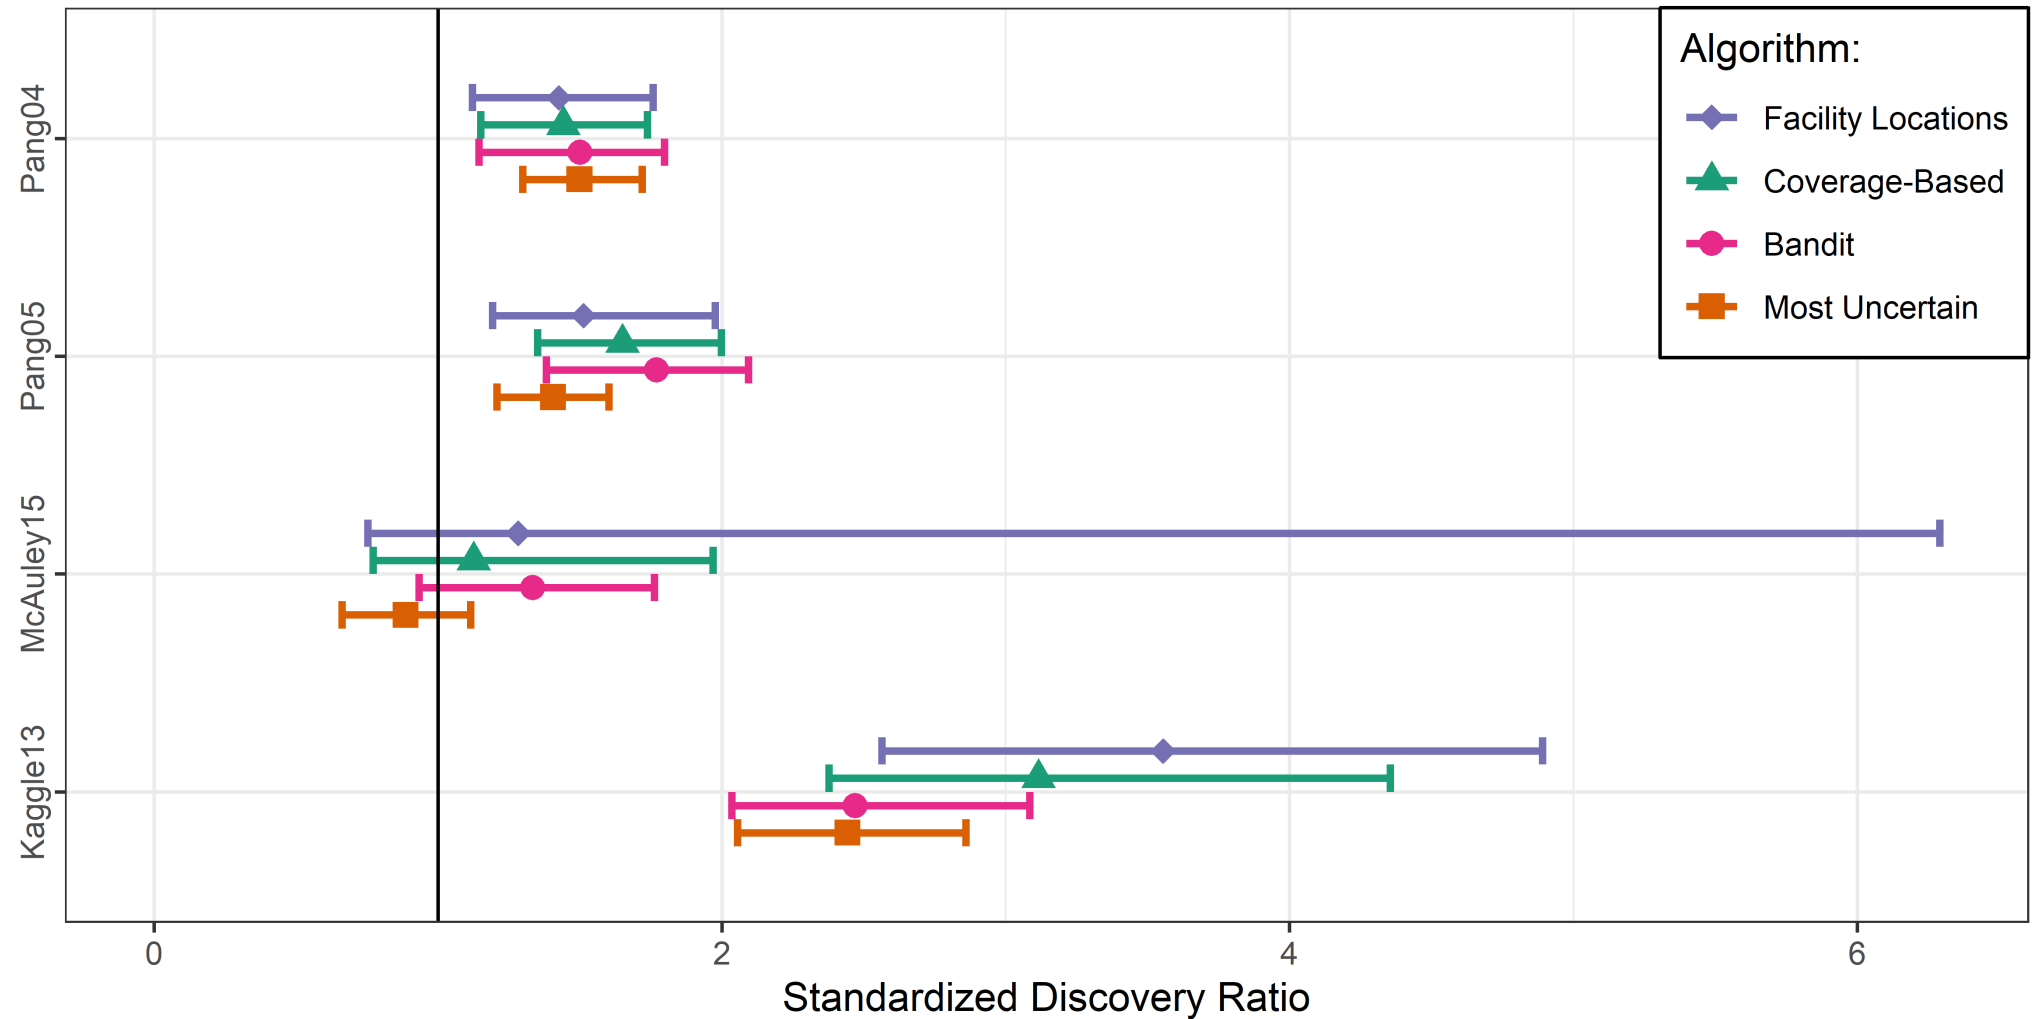
*Standardized Discovery Ratio* (SDR) is the ratio of discovered misclassifications to the expected number based on confidence

Numerator is count of discovered misclassification

Denominator is expected count based on confidence

Higher SDR values suggest higher efficiency of identifying overconfident instances

## Facility Location Query vs. Most-Uncertain Query



# Adversarial Distances for Unsupervised Query Set Selection

# Leveraging Adversarial Learning Methods for Discovering Overconfidence

**Problem:** Adaptive query algorithms appear to discover overconfidence slowly and inconsistently

**Goal:** Initializing oracle query search with selection of "low hanging fruit" - a set of instances in the unlabeled test set that are *likely* to be overconfident

Need unsupervised method to extract more information about model confidence on unlabeled points

*Adversarial instances* are points that have been adapted to fool classifiers with malicious intent

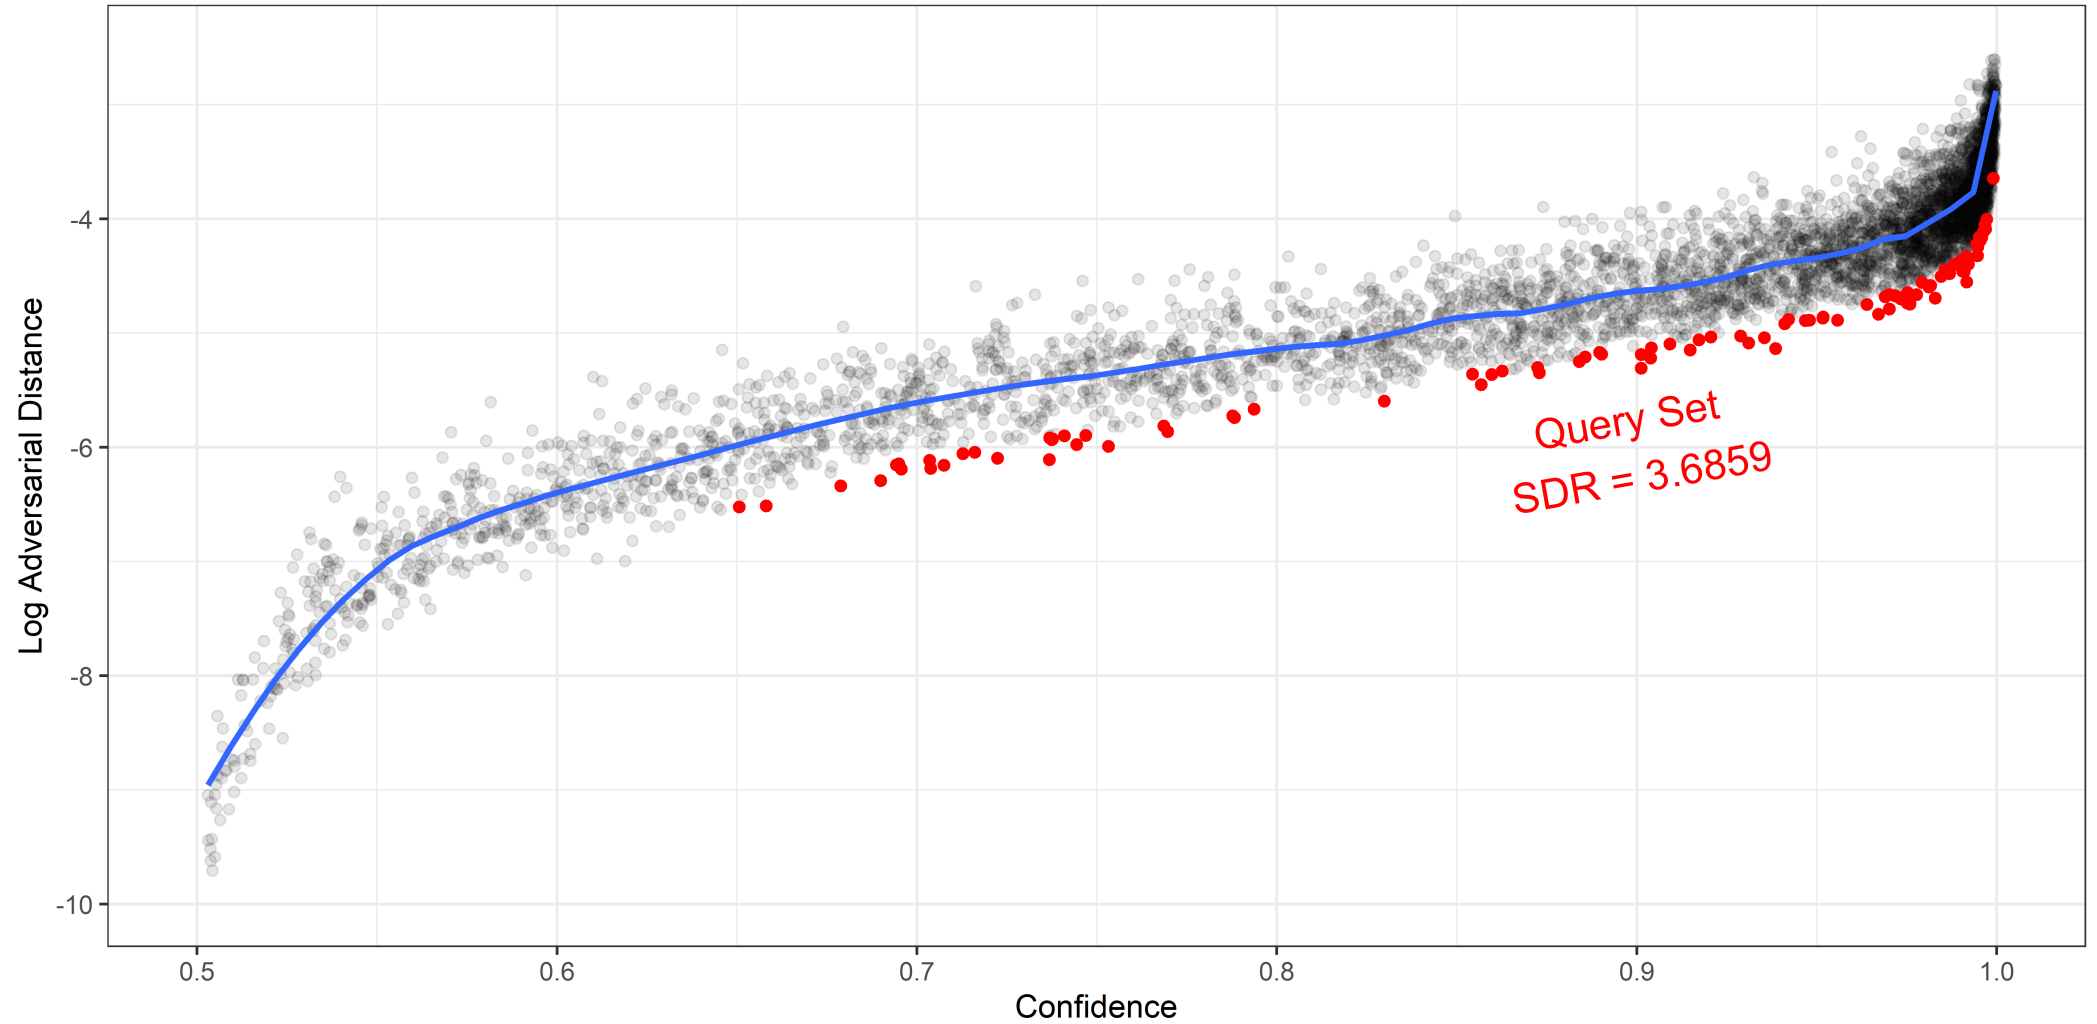
- This can involve minimally changing features until the instance is *misclassified*
- Images can be made adversarial by selectively adding noise in the pixels
- This is a supervised task, because it uses the true class labels

**Proposal:** Use adversarial methods to minimally change features until the instance *changes predicted class*.

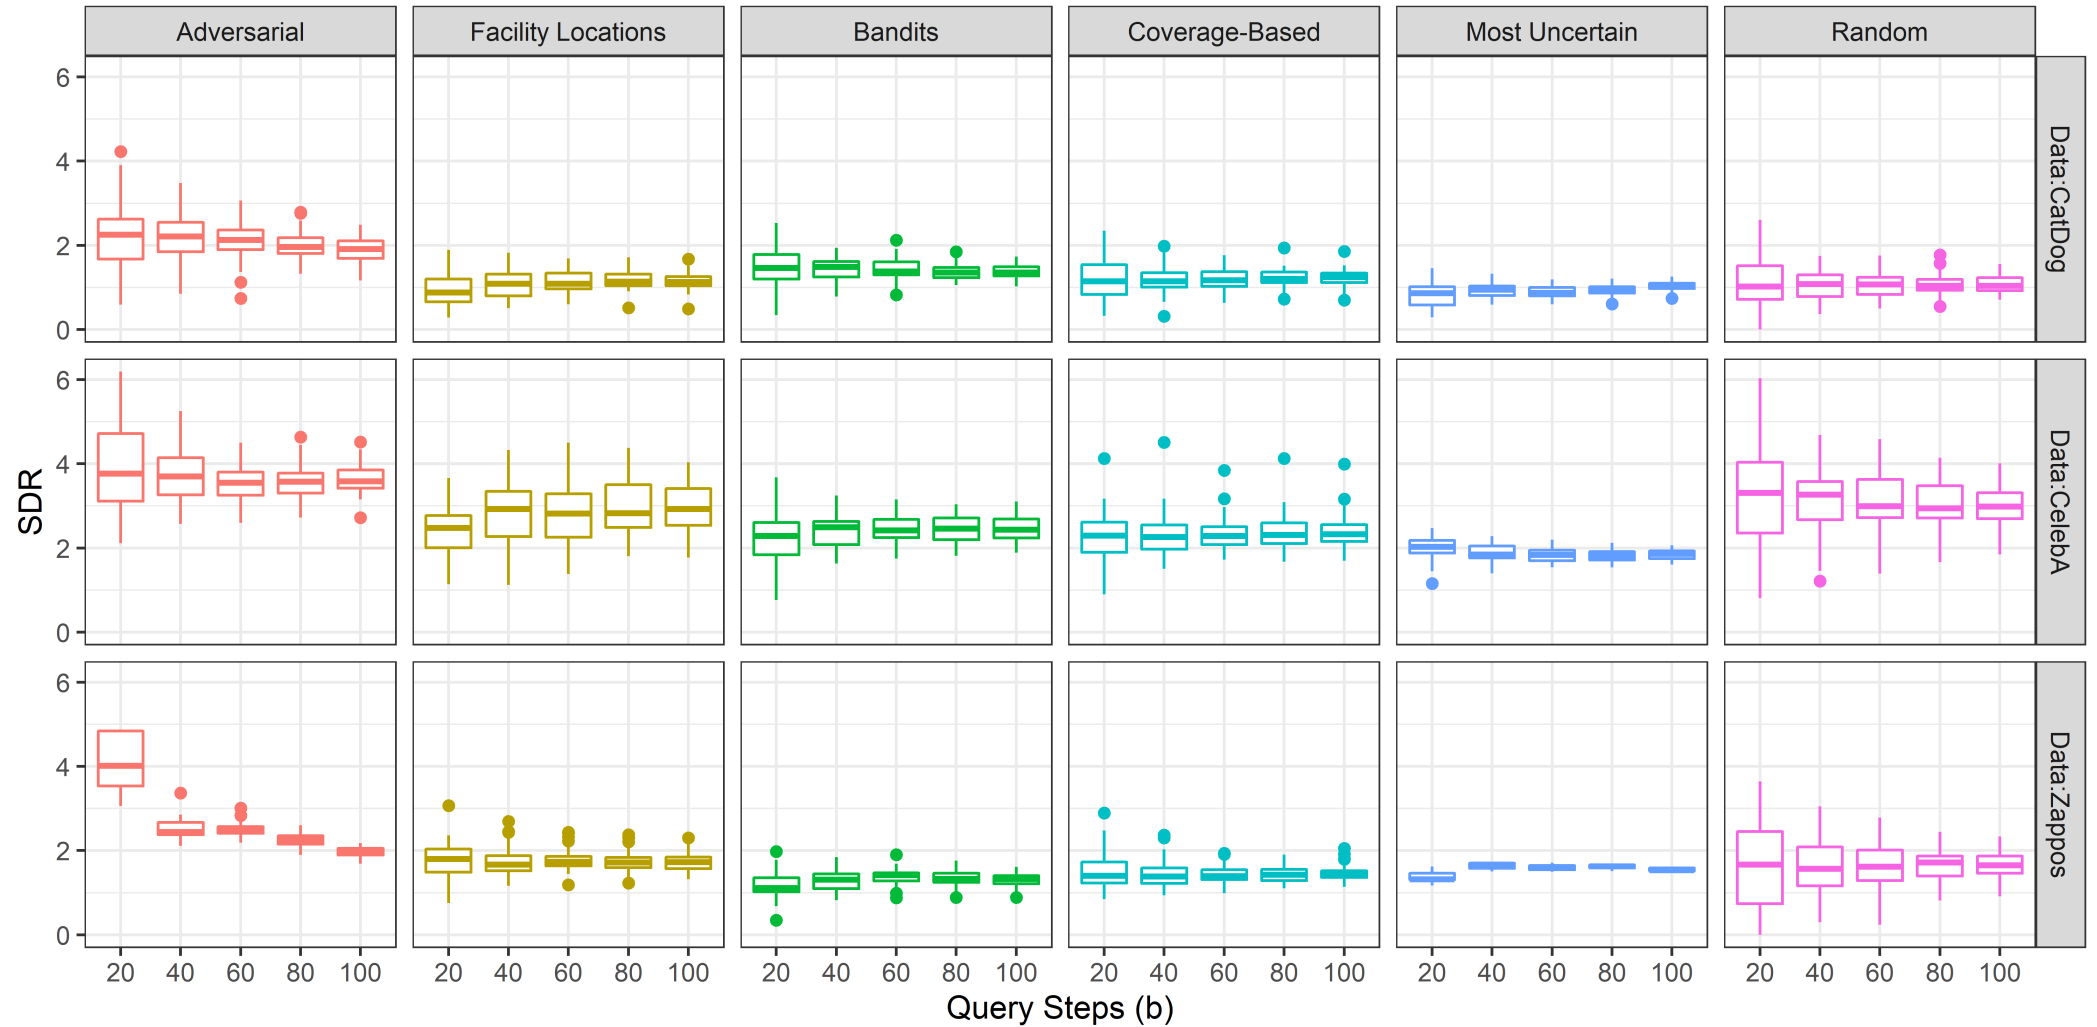
Define *minimal adversarial distance* (MAD) as the distance between an instance and the closest point in the feature space with different predicted class

We will query instances that have low MAD values relative to their confidence score.

## Adversarial Distance-Based Query



# Empirical Experiment Results



# Discussion



# Conclusions

## **Facility Locations Utility Methods**

Performs well across all overconfidence profiles seen in empirical evaluations

SDR improves as budget increases, but this learning can be slow and inconsistent

## **Adversarial Distance Methods**

Efficient in discovering instances that display classifier overconfidence with small budgets

SDR decreases as budget increases, decreasing returns on oracle investment

Implemented successfully in image classifier setting

## Continuing and Future work

### Student Research Projects:

- **Bunyod Tusmatov** extended adversarial methods to text classifiers, observed similarly strong SDR outcomes
- **Sally Dufek** currently working to extend adversarial methods to generalized feature space with black box classifiers
- Can envision **undergraduates** helping evaluate algorithms in a more extensive empirical experiment

### Hybrid query algorithm

- Initialize query with set from adversarial distance method
- Continue with adaptive Facility Location utility method until budget

# Resources

## Code Repository

- These Slides: <https://github.com/kmaurer/search2020/tree/master/ResearchTalk>
- Facility Locations Utility: <https://github.com/kmaurer/uuutils>

## Software and Data

- Garrett, Nar, Fisher, Maurer (2018). ggvoronoi: Voronoi Diagrams and Heatmaps with ggplot2. J. Open Source Software, 3(32), 1096.
- Kuhn (2019). caret: Classification and Regression Training. R package version 6.0-84. <https://CRAN.R-project.org/package=caret>
- R Core Team (2019). R: A language and environment for statistical computing.
- R Foundation for Statistical Computing, Vienna, Austria. URL <https://www.R-project.org/>.
- RStudio Team (2016). RStudio: Integrated Development for R. RStudio, Inc., Boston, MA URL <http://www.rstudio.com/>.
- Wickham (2017). tidyverse: Easily Install and Load the 'Tidyverse'. R package version 1.2.1. <https://CRAN.R-project.org/package=tidyverse>
- Wickham (2019). rvest: Easily Harvest (Scrape) Web Pages. R package version 0.3.4. <https://CRAN.R-project.org/package=rvest>
- Xie (2019). knitr: A General-Purpose Package for Dynamic Report Generation in R. R package version 1.22.
- Zillow home listings data from Oxford, Ohio. <https://www.zillow.com/>

## References

- Arya, V.; Garg, N.; Khandekar, R.; Meyerson, A.; Munagala, K.; and Pandit, V. 2004. Local search heuristics for k-median and facility location problems. *SIAM Journal on computing* 33(3):544–562.
- Attenberg, J.; Ipeirotis, P.; and Provost, F. 2015. Beat the Machine. *Journal of Data and Information Quality* 6(1):1–17.
- Bansal, G., and Weld, D. S. 2018. A coverage-based utility model for identifying unknown unknowns. In *AAAI*.
- Bella, A.; Ferri, C.; Hernández-Orallo, J.; and Ramírez-Quintana, M. J. 2010. Calibration of machine learning models. In *Handbook of Research on Machine Learning Applications and Trends: Algorithms, Methods, and Techniques*. IGI Global. 128–146.
- Brendel, W. ; Rauber, J.; and Bethge, M. 2017. Decision-based adversarial attacks: Reliable attacks against black-box machine learning models. *arXiv preprint arXiv:1712.04248*, 2017. 2, 6
- Casella, G., and Berger, R. L. 2002. *Statistical inference*, volume 2. Duxbury Pacific Grove, CA.
- Farahani, R. Z., and Hekmatfar, M. 2009. *Facility location: concepts, models, algorithms and case studies*. Springer.
- Goodfellow, I. J.; Shlens, J.; and Szegedy, C. 2015. Explaining and Harnessing Adversarial Examples. *Iclr*, pages 1–11, 2.
- Guha, S., and Khuller, S. 1999. Greedy strikes back: Improved facility location algorithms. *Journal of algorithms* 31(1):228–248.
- Guo, C.; Pleiss, G.; Sun, Y.; and Weinberger, K. Q.. On Calibration of Modern Neural Networks. In *International Conference on Machine Learning*, 2017. 3, 5
- Kaggle. 2013. [www.kaggle.com/c/dogs-vs-cats](http://www.kaggle.com/c/dogs-vs-cats). Accessed: 2018-08-28.

## References (continued)

- Lakkaraju, H.; Kamar, E.; Caruana, R.; and Horvitz, E. 2017. Identifying unknown unknowns in the open world: Representations and policies for guided exploration. In AAAI.
- Liu, Z.; Luo, P.; Wang, X.; and Tang X. 2015. Deep learning face attributes in the wild. In Proceedings of International Conference on Computer Vision (ICCV), 5
- McAuley, J.; Pandey, R.; and Leskovec, J. 2015. Inferring networks of substitutable and complementary products. In Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, 785–794. ACM.
- Nushi, B.; Kamar, E.; Horvitz, E.; and Kossmann, D. 2016. On Human Intellect and Machine Failures: Troubleshooting Integrative Machine Learning Systems.
- Ozan, S.; and Savarese, S. 2018. Active Learning for Convolutional Neural Networks: A Core-Set Approach. In ICLR 2018, pages 1–13, 2018. 1
- Pang, B., and Lee, L. 2004. A sentimental education: Sentiment analysis using subjectivity summarization based on minimum cuts. In Proceedings of the 42nd annual meeting on Association for Computational Linguistics, 271. Association for Computational Linguistics.
- Pang, B., and Lee, L. 2005. Seeing stars: Exploiting class relationships for sentiment categorization with respect to rating scales. In Proceedings of the 43rd annual meeting on association for computational linguistics, 115–124. Association for Computational Linguistics.
- Patel, V.; Gopalan, R.; Li, R.; and Chellappa, R. 2014. Visual Domain Adaptation: An Overview of Recent Advances. 1–34.
- Rauber, J.; Brendel, W.; and Bethge, M. 2017. Foolbox: A python toolbox to benchmark the robustness of machine learning models. arXiv preprint arXiv:1707.04131, 2017. 2

## References (continued)

- Ribeiro, M. T.; Singh, S.; and Guestrin, C. 2016. Why Should I Trust You: Explaining the Predictions of Any Classifier. 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, 2016. 6
- Rosner, B. 2015. Fundamentals of biostatistics. Nelson Education.
- Settles, B. 2010. Active Learning Literature Survey. Machine Learning 15(2):201–221.
- Stock, P., and Cisse, M. 2017. Convnets and imagenet beyond accuracy: Explanations, bias detection, adversarial examples and model criticism. arXiv preprint arXiv:1711.11443.
- Sugiyama, M.; Lawrence, N. D.; Schwaighofer, A.; et al. 2017. Dataset shift in machine learning. The MIT Press.
- Taylor, P. 2013. Standardized mortality ratios. International journal of epidemiology 42(6):1882–1890.
- Wang, K.; Zhang, D.; Li, Y.; Zhang, R.; and Lin, L. 2017. CostEffective Active Learning for Deep Image Classification. pages 1–10, 2017. 1
- Yu, A.; and Grauman, K. 2014. Fine-grained visual comparisons with local learning. In Computer Vision and Pattern Recognition (CVPR), Jun 2014. 5
- Yu A.; and Grauman, K. 2017. Semantic jitter: Dense supervision for visual comparisons via synthetic images. In International Conference on Computer Vision (ICCV), Oct 2017. 5
- Zheng, B.; Lin, X.; Xiao, Y.; Yang, J.; and He, L. 2018. An effective method for identifying unknown unknowns with noisy oracle. In International Conference on Case-Based Reasoning, pages 480–495. Springer, 2018. 2

# Thanks!

Slides created via the R package **xaringan**.